

A COMPLETE PRACTICAL GUIDE TO MASTER
DEEP LEARNING CONCEPTS AND BUILD REAL-WORLD AI SOLUTIONS



DEEP LEARNING FUNDAMENTALS

— FROM ZERO TO HERO —



Build Strong Foundations



Hands-on Deep Learning with Python




Neural Networks to Advanced Architectures



Real-World Projects & Production Deployment



MLOps, Best Practices & Career Insights



COVERING EVERYTHING YOU NEED TO KNOW



CORE
CONCEPTS



NEURAL
NETWORKS



CNNs




RNNs &
LSTMs



TRANSFORMERS
& ATTENTION



GANs &
GENERATIVE AI



TRANSFER
LEARNING



MLOps &
DEPLOYMENT



AZURE AI
INTEGRATION



REAL-WORLD
PROJECTS

BY ABHISHEK PANDA

— AI ARCHITECT | AZURE AI ARCHITECT —



BEGINNER TO
ADVANCED



PRACTICAL CODE
EXAMPLES



BEST PRACTICES
& TIPS



CAREER
ACCELERATION

DEEP LEARNING FUNDAMENTALS: FROM ZERO TO HERO

BY ABHISHEK PANDA
AI ARCHITECT | AZURE AI ARCHITECT

PART I — DEEP LEARNING FOUNDATIONS

CHAPTER 1: INTRODUCTION TO DEEP LEARNING

- 1.1 WHAT IS DEEP LEARNING?
- 1.2 AI VS ML VS DEEP LEARNING VS GENERATIVE AI
- 1.3 HISTORY AND EVOLUTION OF NEURAL NETWORKS
- 1.4 WHY DEEP LEARNING WORKS
- 1.5 REPRESENTATION LEARNING EXPLAINED
- 1.6 REAL-WORLD DEEP LEARNING APPLICATIONS
- 1.7 DEEP LEARNING LIFECYCLE
- 1.8 HARDWARE REQUIREMENTS (CPU, GPU, TPU, NPU, LPU, DPU)
- 1.9 DEEP LEARNING CAREER PATHS

CHAPTER 2: MACHINE LEARNING BRIDGE FOR DEEP LEARNING

- 2.1 WHY MACHINE LEARNING MATTERS BEFORE DEEP LEARNING
- 2.2 CLASSICAL ML VS DEEP LEARNING
- 2.3 FEATURE ENGINEERING VS REPRESENTATION LEARNING
- 2.4 WHEN TO USE ML VS DL
- 2.5 DEEP LEARNING PROBLEM FRAMING
- 2.6 COMMON DL USE CASES BY DOMAIN

CHAPTER 3: MATHEMATICS FOR DEEP LEARNING

- 3.1 LINEAR ALGEBRA FOR NEURAL NETWORKS
- 3.2 SCALARS, VECTORS, MATRICES, TENSORS
- 3.3 MATRIX MULTIPLICATION IN NEURAL NETWORKS
- 3.4 EIGENVALUES AND EIGENVECTORS
- 3.5 VECTOR CALCULUS
- 3.6 DERIVATIVES AND PARTIAL DERIVATIVES
- 3.7 CHAIN RULE
- 3.8 GRADIENTS AND JACOBIANS
- 3.9 HESSIANS AND SECOND-ORDER DERIVATIVES
- 3.10 PROBABILITY THEORY FOR DL
- 3.11 STATISTICS FOR MODEL TRAINING

DEEP LEARNING ENVIRONMENT SETUP

- 3.1 PYTHON ENVIRONMENT SETUP
- 3.2 CUDA / GPU SETUP
- 3.3 JUPYTER / VS CODE WORKFLOW
- 3.4 VIRTUAL ENVIRONMENTS
- 3.5 WORKING WITH LARGE DATASETS
- 3.6 EXPERIMENT REPRODUCIBILITY

CHAPTER 4: PROGRAMMING FOUNDATIONS

- 4.1 PYTHON FOR DEEP LEARNING
- 4.2 NUMPY FOR TENSOR OPERATIONS
- 4.3 PANDAS FOR DATA PROCESSING
- 4.4 DATA VISUALIZATION BASICS
- 4.5 GIT/GITHUB FOR EXPERIMENT VERSIONING
- 4.6 SQL FOR DEEP LEARNING DATA PIPELINES

PART II — DEEP LEARNING FRAMEWORKS & TOOLING

CHAPTER 5: TENSORFLOW ECOSYSTEM

- 5.1 TENSORFLOW FUNDAMENTALS
- 5.2 TENSOR OPERATIONS
- 5.3 COMPUTATIONAL GRAPHS
- 5.4 EAGER EXECUTION
- 5.5 TENSORFLOW BEST PRACTICES

CHAPTER 6: KERAS DEEP DIVE

- 6.1 KERAS SEQUENTIAL API
- 6.2 FUNCTIONAL API
- 6.3 MODEL SUBCLASSING
- 6.4 CUSTOM LAYERS
- 6.5 CUSTOM TRAINING LOOPS

CHAPTER 7: PYTORCH DEEP DIVE

- 7.1 TENSORS AND AUTOGRAD
- 7.2 NN.MODULE ARCHITECTURE
- 7.3 DATASET AND DATALOADER
- 7.4 BUILDING TRAINING LOOPS
- 7.5 CUSTOM MODULES AND LAYERS

CHAPTER 8: FASTAI & RAPID PROTOTYPING

- 8.1 FASTAI FUNDAMENTALS
- 8.2 DATABLOCK API
- 8.3 LEARNER API
- 8.4 TRANSFER LEARNING WITH FASTAI

PART III — NEURAL NETWORK FUNDAMENTALS

CHAPTER 9: ARTIFICIAL NEURONS AND PERCEPTRONS

- 9.1 BIOLOGICAL NEURON VS ARTIFICIAL NEURON
- 9.2 PERCEPTRON MODEL
- 9.3 ACTIVATION FUNCTIONS OVERVIEW
- 9.4 LIMITATIONS OF SINGLE LAYER PERCEPTRONS

CHAPTER 10: ARTIFICIAL NEURAL NETWORKS (ANNs)

- 10.1 FEEDFORWARD NEURAL NETWORKS
- 10.2 FORWARD PROPAGATION
- 10.3 LOSS FUNCTIONS
- 10.4 BACKPROPAGATION
- 10.5 GRADIENT DESCENT
- 10.6 WEIGHT INITIALIZATION
- 10.7 UNIVERSAL APPROXIMATION THEOREM
- 10.8 BUILDING ANN MODELS FROM SCRATCH

CHAPTER 11: ACTIVATION FUNCTIONS

- 11.1 SIGMOID
- 11.2 TANH
- 11.3 RELU
- 11.4 LEAKY RELU
- 11.5 ELU
- 11.6 GELU
- 11.7 SWISH
- 11.8 SOFTMAX

CHAPTER 12: LOSS FUNCTIONS

- 12.1 MSE
- 12.2 MAE
- 12.3 BINARY CROSS ENTROPY
- 12.4 CATEGORICAL CROSS ENTROPY
- 12.5 SPARSE CATEGORICAL CROSS ENTROPY
- 12.6 HINGE LOSS
- 12.7 CONTRASTIVE LOSS
- 12.8 TRIPLET LOSS
- 12.9 FOCAL LOSS

PART IV — OPTIMIZATION & REGULARIZATION

CHAPTER 13: OPTIMIZATION ALGORITHMS

- 13.1 BATCH GRADIENT DESCENT
- 13.2 STOCHASTIC GRADIENT DESCENT (SGD)
- 13.3 MINI-BATCH GRADIENT DESCENT
- 13.4 MOMENTUM
- 13.5 NESTEROV ACCELERATED GRADIENT
- 13.6 ADAGRAD
- 13.7 RMSPROP
- 13.8 ADAM
- 13.9 ADAMW
- 13.10 LION OPTIMIZER

CHAPTER 14: TRAINING BEST PRACTICES

- 14.1 LEARNING RATE SCHEDULING
- 14.2 WARMUP STRATEGIES
- 14.3 GRADIENT CLIPPING
- 14.4 MIXED PRECISION TRAINING
- 14.5 DISTRIBUTED TRAINING
- 14.6 TRANSFER LEARNING FUNDAMENTALS
- 14.7 FINE TUNING STRATEGIES

CHAPTER 15: REGULARIZATION TECHNIQUES

- 15.1 L1 / L2 REGULARIZATION
- 15.2 DROPOUT
- 15.3 BATCH NORMALIZATION
- 15.4 LAYER NORMALIZATION
- 15.5 DATA AUGMENTATION
- 15.6 EARLY STOPPING
- 15.7 WEIGHT DECAY
- 15.8 LABEL SMOOTHING

PART V — COMPUTER VISION DEEP LEARNING

CHAPTER 16: CONVOLUTIONAL NEURAL NETWORKS [CNN]

- 16.1 CONVOLUTION OPERATION
- 16.2 FILTERS AND KERNELS
- 16.3 PADDING AND STRIDE
- 16.4 POOLING LAYERS
- 16.5 FEATURE MAPS
- 16.6 RECEPTIVE FIELDS

CHAPTER 17: CLASSIC CNN ARCHITECTURES

- 17.1 LENET
- 17.2 ALEXNET
- 17.3 VGGNET
- 17.4 GOOGLNET / INCEPTION
- 17.5 RESNET
- 17.6 DENSENET
- 17.7 EFFICIENTNET
- 17.8 CONVNEXT

CHAPTER 18: ADVANCED COMPUTER VISION TASKS

- 18.1 IMAGE CLASSIFICATION
- 18.2 OBJECT DETECTION FUNDAMENTALS
- 18.3 RCNN FAMILY
- 18.4 SSD
- 18.5 YOLO FAMILY
- 18.6 RETINANET
- 18.7 SEMANTIC SEGMENTATION
- 18.8 U-NET
- 18.9 MASK R-CNN
- 18.10 IMAGE SIMILARITY / EMBEDDINGS

PART VI — SEQUENTIAL & TEMPORAL DEEP LEARNING

CHAPTER 19: RECURRENT NEURAL NETWORKS [RNN]

- 19.1 SEQUENTIAL DATA FUNDAMENTALS
- 19.2 VANILLA RNN
- 19.3 BACKPROPAGATION THROUGH TIME (BPTT)
- 19.4 VANISHING / EXPLODING GRADIENTS

CHAPTER 20: ADVANCED SEQUENCE ARCHITECTURES

- 20.1 LSTM
- 20.2 GRU
- 20.3 BIDIRECTIONAL RNNs
- 20.4 SEQ2SEQ MODELS
- 20.5 ENCODER-DECODER ARCHITECTURE

PART VII — ATTENTION & TRANSFORMERS

CHAPTER 21: ATTENTION MECHANISMS

- 21.1 WHY ATTENTION MATTERS
- 21.2 ADDITIVE VS MULTIPLICATIVE ATTENTION
- 21.3 SCALED DOT PRODUCT ATTENTION
- 21.4 SELF ATTENTION
- 21.5 CROSS ATTENTION
- 21.6 MULTI-HEAD ATTENTION

CHAPTER 22: TRANSFORMER ARCHITECTURE

- 22.1 ENCODER ARCHITECTURE
- 22.2 DECODER ARCHITECTURE
- 22.3 POSITIONAL ENCODING
- 22.4 RESIDUAL CONNECTIONS
- 22.5 LAYER NORMALIZATION IN TRANSFORMERS

CHAPTER 23: TRANSFORMER VARIANTS

- 23.1 BERT
- 23.2 ROBERTA
- 23.3 GPT FAMILY
- 23.4 T5
- 23.5 VISION TRANSFORMERS
- 23.6 SWIN TRANSFORMERS
- 23.7 TIME SERIES TRANSFORMERS

PART VIII — GENERATIVE DEEP LEARNING

CHAPTER 24: AUTOENCODERS

- 24.1 ENCODER/DECODER ARCHITECTURE
- 24.2 SPARSE AUTOENCODERS
- 24.3 DENOISING AUTOENCODERS
- 24.4 VARIATIONAL AUTOENCODERS

CHAPTER 25: GENERATIVE ADVERSARIAL NETWORKS [GANs]

- 25.1 GAN FUNDAMENTALS
- 25.2 GENERATOR VS DISCRIMINATOR
- 25.3 DCGAN
- 25.4 CONDITIONAL GAN

25.5 CYCLEGAN
25.6 STYLEGAN

CHAPTER 26: DIFFUSION MODELS

26.1 DIFFUSION FUNDAMENTALS
26.2 FORWARD / REVERSE PROCESS
26.3 DDPM
26.4 STABLE DIFFUSION CONCEPTS
26.5 LATENT DIFFUSION MODELS

PART IX — MODERN DEEP LEARNING ARCHITECTURES

CHAPTER 27: FOUNDATION MODELS

27.1 WHAT ARE FOUNDATION MODELS
27.2 PRETRAINING PARADIGMS
27.3 TRANSFER LEARNING AT SCALE
27.4 FINE-TUNING VS PROMPTING
27.5 PARAMETER EFFICIENT FINE-TUNING (PEFT)
27.6 LOW-RANK ADAPTATION (LORA)
27.7 ADAPTER LAYERS
27.8 MODEL DISTILLATION

CHAPTER 28: MULTIMODAL DEEP LEARNING

28.1 VISION-LANGUAGE MODELS
28.2 CLIP
28.3 BLIP
28.4 AUDIO-VISUAL LEARNING
28.5 CROSS-MODAL EMBEDDINGS

CHAPTER 29: GRAPH NEURAL NETWORKS

29.1 GRAPH FUNDAMENTALS
29.2 GRAPH CONVOLUTIONAL NETWORKS
29.3 GRAPH ATTENTION NETWORKS
29.4 GRAPHSAGE
29.5 GNN USE CASES

PART X — DEEP REINFORCEMENT LEARNING

CHAPTER 30: DEEP RL FUNDAMENTALS

- 30.0 RL RECAP
- 30.1 DQN
- 30.2 DOUBLE DQN
- 30.3 DUELING DQN
- 30.4 PPO
- 30.5 SAC
- 30.6 MULTI-AGENT RL OVERVIEW

PART XI — EXPLAINABILITY, SAFETY, AND ROBUSTNESS

CHAPTER 31: EXPLAINABLE DEEP LEARNING

- 31.1 SALIENCY MAPS
- 31.2 GRAD-CAM
- 31.3 SHAP FOR DL
- 31.4 INTEGRATED GRADIENTS

CHAPTER 32: SAFETY & ROBUSTNESS

- 32.1 ADVERSARIAL ATTACKS
- 32.2 ADVERSARIAL TRAINING
- 32.3 ROBUSTNESS TESTING
- 32.4 BIAS AND FAIRNESS IN DL

PART XII — PRODUCTION DEEP LEARNING

CHAPTER 33: DEPLOYMENT FUNDAMENTALS

- 33.1 MODEL SERIALIZATION
- 33.2 FASTAPI DEPLOYMENT
- 33.3 STREAMLIT / GRADIO INTERFACES
- 33.4 REST / GRPC INFERENCE APIS

CHAPTER 34: MODEL OPTIMIZATION FOR INFERENCE

- 34.1 QUANTIZATION
- 34.2 PRUNING
- 34.3 DISTILLATION

34.4 ONNX CONVERSION
34.5 TENSORRT OPTIMIZATION

CHAPTER 35: PRODUCTION SERVING INFRASTRUCTURE

35.1 DOCKER FOR DEEP LEARNING
35.2 KUBERNETES FOR GPU WORKLOADS
35.3 NVIDIA TRITON INFERENCE SERVER
35.4 BATCH VS REAL-TIME INFERENCE

PART XIII — AZURE DEEP LEARNING & MLOPS

CHAPTER 36: AZURE DEEP LEARNING PLATFORM

36.1 AZURE ML WORKSPACE
36.2 GPU COMPUTE CLUSTERS
36.3 DISTRIBUTED TRAINING
36.4 HYPERPARAMETER SWEEPS
36.5 EXPERIMENT TRACKING

CHAPTER 37: AZURE DEPLOYMENT & SERVING

37.1 MANAGED ONLINE ENDPOINTS
37.2 BATCH ENDPOINTS
37.3 AKS GPU DEPLOYMENT
37.4 AZURE CONTAINER APPS

CHAPTER 38: MLOps & LLMOPS FUNDAMENTALS

38.1 MLFLOW / W&B / TENSORBOARD
38.2 CI/CD FOR DEEP LEARNING

PART XIV — DEEP LEARNING MLOPS

CHAPTER 39: DL EXPERIMENT TRACKING

39.1 MLFLOW
39.2 WEIGHTS & BIASES
39.3 TENSORBOARD TRACKING

CHAPTER 40: MONITORING DEEP LEARNING SYSTEMS

40.1 DRIFT DETECTION

40.2 EMBEDDING MONITORING

40.3 LATENCY / THROUGHPUT METRICS

40.4 RETRAINING PIPELINES

PART XV — ENTERPRISE PROJECTS

CHAPTER 41: BEGINNER PROJECTS

1. MNIST HANDWRITTEN DIGIT RECOGNITION [(INTRO TO ANN/CNN)]
2. CIFAR-10 MULTI-CLASS IMAGE CLASSIFICATION [(DEEPER CNN / DATA AUGMENTATION)]

CHAPTER 42: INTERMEDIATE PROJECTS

1. TRANSFER LEARNING IMAGE CLASSIFIER WITH RESNET/EFFICIENTNET
(REAL-WORLD TRANSFER LEARNING)
2. SENTIMENT ANALYSIS WITH LSTM / GRU
(NLP SEQUENCE MODELING)
3. TIME SERIES FORECASTING WITH LSTM
(SEQUENTIAL REGRESSION)

CHAPTER 43: ADVANCED PROJECTS

1. OBJECT DETECTION PRODUCTION PIPELINE [(YOLO/SSD)]
2. GAN IMAGE GENERATOR
3. VISION TRANSFORMER PIPELINE
4. MULTIMODAL VISION-LANGUAGE PIPELINE
5. AZURE END-TO-END DL PRODUCTION SYSTEM
6. ENTERPRISE MLOPS FOR DEEP LEARNING

APPENDICES

APPENDIX A — RESEARCH PAPER READING FRAMEWORK

APPENDIX B — CUDA / GPU OPTIMIZATION GUIDE

APPENDIX C — DEEP LEARNING INTERVIEW PREPARATION

APPENDIX D — AZURE CERTIFICATION ROADMAP