

A COMPLETE PRACTICAL GUIDE TO MASTER
NATURAL LANGUAGE PROCESSING AND BUILD REAL-WORLD AI SOLUTIONS



NLP MASTERY

— ZERO TO HERO —

FROM FUNDAMENTALS TO TRANSFORMERS
AND BEYOND



Covers Everything: Basics to Advanced NLP, Transformers, LLM Foundations & Production



Text Processing, Embeddings, Models, and Real-World Applications



Hands-on with Transformers, BERT, GPT, LangChain Essentials & RAG Concepts



Deployment, MLOps, and Scalable NLP Systems with Azure



Industry-Ready Projects, Best Practices & Interview Preparation



PRACTICAL
EXAMPLES



REAL-WORLD
PROJECTS



SOTA MODELS
& TECHNIQUES



AZURE NLP
INTEGRATION



MLOps &
DEPLOYMENT

MASTER CORE NLP TASKS



TEXT
CLASSIFICATION



SENTIMENT
ANALYSIS



NAMED ENTITY
RECOGNITION



MACHINE
TRANSLATION



TEXT
SUMMARIZATION



QUESTION
ANSWERING





CHATBOTS &
CONVERSATIONAL AI



TEXT
GENERATION

POWERED BY



BY ABHISHEK PANDA

AI ARCHITECT | AZURE AI ARCHITECT



LEARN SMARTER
BUILD FASTER



PRACTICAL.
REAL. IMPACTFUL.



BEGINNER FRIENDLY
EXPERT FOCUSED



FUTURE-READY
NLP SKILLS

NATURAL LANGUAGE PROCESSING: FROM ZERO TO HERO

BY ABHISHEK PANDA
AI ARCHITECT | AZURE AI ARCHITECT

PART I — FOUNDATIONS OF NATURAL LANGUAGE PROCESSING

CHAPTER 1: INTRODUCTION TO NLP

- 1.1 WHAT IS NATURAL LANGUAGE PROCESSING?
- 1.2 WHY NLP MATTERS IN MODERN AI
- 1.3 NLP VS NLU VS NLG
- 1.4 HISTORY AND EVOLUTION OF NLP
- 1.5 RULE-BASED VS STATISTICAL VS NEURAL NLP
- 1.6 REAL-WORLD APPLICATIONS OF NLP
- 1.7 NLP LIFECYCLE OVERVIEW
- 1.8 CHALLENGES IN NATURAL LANGUAGE UNDERSTANDING

CHAPTER 2: LINGUISTICS FOUNDATIONS FOR NLP

- 2.1 MORPHOLOGY
- 2.2 SYNTAX
- 2.3 SEMANTICS
- 2.4 PRAGMATICS
- 2.5 DISCOURSE ANALYSIS
- 2.6 AMBIGUITY IN LANGUAGE
- 2.7 CONTEXT IN HUMAN LANGUAGE

CHAPTER 3: PHASES OF NLP PIPELINE

- 3.1 TEXT COLLECTION
- 3.2 DATA CLEANING
- 3.3 TEXT PREPROCESSING
- 3.4 FEATURE EXTRACTION
- 3.5 MODELING
- 3.6 EVALUATION
- 3.7 DEPLOYMENT

PART II — NLP DEVELOPMENT ENVIRONMENT & LIBRARIES

CHAPTER 4: NLP LIBRARIES AND ECOSYSTEM

- 4.1 NLTK
- 4.2 SPACY
- 4.3 TEXTBLOB
- 4.4 GENSIM
- 4.5 HUGGING FACE TRANSFORMERS
- 4.6 DATASETS LIBRARY
- 4.7 TOKENIZERS LIBRARY
- 4.8 SENTENCE TRANSFORMERS

CHAPTER 5: PROGRAMMING FOUNDATIONS FOR NLP

- 5.1 PYTHON FOR NLP
- 5.2 REGEX FOR TEXT PROCESSING
- 5.3 FILE HANDLING AND CORPUS PROCESSING
- 5.4 WORKING WITH LARGE TEXT DATASETS
- 5.5 SQL / NOSQL FOR NLP PIPELINES

PART III — TEXT PREPROCESSING & LINGUISTIC PROCESSING

CHAPTER 6: TEXT CLEANING AND NORMALIZATION

- 6.1 LOWERCASING AND CASING STRATEGIES
- 6.2 PUNCTUATION REMOVAL
- 6.3 STOPWORD REMOVAL
- 6.4 WHITESPACE NORMALIZATION
- 6.5 UNICODE AND ENCODING HANDLING
- 6.6 SPECIAL CHARACTER CLEANING
- 6.7 TEXT NORMALIZATION TECHNIQUES

CHAPTER 7: TOKENIZATION TECHNIQUES

- 7.1 WORD TOKENIZATION
- 7.2 SENTENCE TOKENIZATION
- 7.3 SUBWORD TOKENIZATION
- 7.4 BYTE PAIR ENCODING (BPE)
- 7.5 WORDPIECE
- 7.6 SENTENCEPIECE

CHAPTER 8: STEMMING AND LEMMATIZATION

- 8.1 PORTER STEMMER
- 8.2 SNOWBALL STEMMER
- 8.3 RULE-BASED LEMMATIZATION
- 8.4 CONTEXT-AWARE LEMMATIZATION

CHAPTER 9: LINGUISTIC FEATURE EXTRACTION

- 9.1 PART-OF-SPEECH TAGGING
- 9.2 CHUNKING
- 9.3 DEPENDENCY PARSING
- 9.4 CONSTITUENCY PARSING
- 9.5 NAMED ENTITY RECOGNITION BASICS
- 9.6 COREFERENCE RESOLUTION

PART IV — CLASSICAL NLP REPRESENTATIONS

CHAPTER 10: TEXT REPRESENTATION FUNDAMENTALS

- 10.1 WHY TEXT MUST BE VECTORIZED
- 10.2 SPARSE VS DENSE REPRESENTATIONS
- 10.3 VOCABULARY CONSTRUCTION

CHAPTER 11: TRADITIONAL TEXT VECTORIZATION

- 11.1 ONE-HOT ENCODING
- 11.2 BAG OF WORDS
- 11.3 N-GRAM MODELS
- 11.4 TF-IDF
- 11.5 FEATURE HASHING

CHAPTER 12: STATISTICAL LANGUAGE MODELING

- 12.1 PROBABILITY OF LANGUAGE
- 12.2 MARKOV ASSUMPTION
- 12.3 N-GRAM LANGUAGE MODELS
- 12.4 SMOOTHING TECHNIQUES
- 12.5 PERPLEXITY

CHAPTER 13: TOPIC MODELING

- 13.1 LATENT SEMANTIC ANALYSIS (LSA)
- 13.2 SINGULAR VALUE DECOMPOSITION (SVD)
- 13.3 LATENT DIRICHLET ALLOCATION (LDA)
- 13.4 NON-NEGATIVE MATRIX FACTORIZATION (NMF)

PART V — WORD EMBEDDINGS & SEMANTIC REPRESENTATIONS

CHAPTER 14: DISTRIBUTED WORD REPRESENTATIONS

- 14.1 WHY DENSE EMBEDDINGS MATTER
- 14.2 WORD2VEC ARCHITECTURE
- 14.3 CBOW VS SKIP-GRAM
- 14.4 NEGATIVE SAMPLING
- 14.5 HIERARCHICAL SOFTMAX

CHAPTER 15: ADVANCED EMBEDDING TECHNIQUES

- 15.1 GLOVE
- 15.2 FASTTEXT
- 15.3 DOC2VEC
- 15.4 SENTENCE EMBEDDINGS

CHAPTER 16: CONTEXTUAL EMBEDDINGS

- 16.1 WHY STATIC EMBEDDINGS FALL SHORT
- 16.2 ELMO
- 16.3 CONTEXTUAL REPRESENTATION LEARNING
- 16.4 BERT EMBEDDINGS
- 16.5 ROBERTA EMBEDDINGS
- 16.6 DISTILBERT EMBEDDINGS

PART VI — TRADITIONAL NLP MACHINE LEARNING

CHAPTER 17: CLASSICAL ML FOR NLP

- 17.1 NAIVE BAYES FOR TEXT
- 17.2 LOGISTIC REGRESSION FOR NLP
- 17.3 SUPPORT VECTOR MACHINES
- 17.4 RANDOM FOREST FOR NLP
- 17.5 FEATURE ENGINEERING FOR TEXT MODELS

CHAPTER 18: EVALUATION METRICS FOR NLP

- 18.1 ACCURACY / PRECISION / RECALL / F1
- 18.2 BLEU SCORE
- 18.3 ROUGE SCORE
- 18.4 METEOR

18.5 PERPLEXITY
18.6 SEMANTIC SIMILARITY METRICS

PART VII — DEEP LEARNING FOR NLP

CHAPTER 19: NEURAL NLP FOUNDATIONS

19.1 ANN FOR NLP
19.2 EMBEDDING LAYERS
19.3 SEQUENCE MODELING FUNDAMENTALS

CHAPTER 20: RECURRENT NEURAL NETWORKS

20.1 VANILLA RNN
20.2 BACKPROPAGATION THROUGH TIME
20.3 VANISHING GRADIENT PROBLEM

CHAPTER 21: ADVANCED SEQUENCE MODELS

21.1 LSTM
21.2 GRU
21.3 BIDIRECTIONAL RNNs
21.4 SEQ2SEQ ARCHITECTURE
21.5 ENCODER-DECODER MODELS

CHAPTER 22: CNNs FOR NLP

22.1 CONVOLUTION FOR TEXT
22.2 TEXTCNN
22.3 CNN FOR SENTENCE CLASSIFICATION

PART VIII — ATTENTION & TRANSFORMERS

CHAPTER 23: ATTENTION MECHANISMS

23.1 WHY ATTENTION CHANGED NLP
23.2 ADDITIVE ATTENTION
23.3 MULTIPLICATIVE ATTENTION
23.4 SELF-ATTENTION
23.5 MULTI-HEAD ATTENTION

CHAPTER 24: TRANSFORMER ARCHITECTURE

- 24.1 ENCODER ARCHITECTURE
- 24.2 DECODER ARCHITECTURE
- 24.3 POSITIONAL ENCODING
- 24.4 RESIDUAL CONNECTIONS
- 24.5 LAYER NORMALIZATION

CHAPTER 25: TRANSFORMER-BASED MODELS

- 25.1 BERT
- 25.2 ROBERTA
- 25.3 DISTILBERT
- 25.4 ALBERT
- 25.5 XLNET
- 25.6 TRANSFORMER XL

CHAPTER 26: GENERATIVE TRANSFORMER MODELS

- 26.1 GPT FAMILY
- 26.2 T5
- 26.3 BART
- 26.4 FLAN-T5
- 26.5 INSTRUCTION-TUNED MODELS

PART IX — ADVANCED NLP & MODERN LANGUAGE MODELS

CHAPTER 27: TRANSFER LEARNING IN NLP

- 27.1 FINE-TUNING PRETRAINED MODELS
- 27.2 DOMAIN ADAPTATION
- 27.3 FEW-SHOT LEARNING
- 27.4 ZERO-SHOT LEARNING

CHAPTER 28: PARAMETER EFFICIENT FINE-TUNING

- 28.1 ADAPTER LAYERS
- 28.2 PREFIX TUNING
- 28.3 PROMPT TUNING
- 28.4 LORA
- 28.5 QLoRA

CHAPTER 29: RETRIEVAL & SEMANTIC SEARCH FOUNDATIONS

- 29.1 SENTENCE TRANSFORMERS
- 29.2 EMBEDDING MODELS
- 29.3 SEMANTIC SIMILARITY
- 29.4 DENSE RETRIEVAL
- 29.5 HYBRID SEARCH CONCEPTS

PART X — CORE NLP TASKS

CHAPTER 30: TEXT CLASSIFICATION

- 30.1 SPAM DETECTION
- 30.2 TOPIC CLASSIFICATION
- 30.3 INTENT CLASSIFICATION
- 30.4 MULTI-LABEL CLASSIFICATION

CHAPTER 31: SENTIMENT ANALYSIS

- 31.1 RULE-BASED SENTIMENT ANALYSIS
- 31.2 VADER
- 31.3 DEEP LEARNING SENTIMENT MODELS
- 31.4 ASPECT-BASED SENTIMENT ANALYSIS

CHAPTER 32: NAMED ENTITY RECOGNITION

- 32.1 NER FUNDAMENTALS
- 32.2 RULE-BASED NER
- 32.3 ML-BASED NER
- 32.4 TRANSFORMER-BASED NER

CHAPTER 33: INFORMATION EXTRACTION

- 33.1 RELATION EXTRACTION
- 33.2 ENTITY LINKING
- 33.3 COREFERENCE RESOLUTION
- 33.4 KNOWLEDGE GRAPH CONSTRUCTION

CHAPTER 34: MACHINE TRANSLATION

- 34.1 STATISTICAL MACHINE TRANSLATION
- 34.2 NEURAL MACHINE TRANSLATION
- 34.3 TRANSFORMER-BASED TRANSLATION

CHAPTER 35: TEXT SUMMARIZATION

- 35.1 EXTRACTIVE SUMMARIZATION**
- 35.2 ABSTRACTIVE SUMMARIZATION**
- 35.3 TRANSFORMER SUMMARIZATION MODELS**

CHAPTER 36: QUESTION ANSWERING SYSTEMS

- 36.1 RETRIEVAL-BASED QA**
- 36.2 EXTRACTIVE QA**
- 36.3 GENERATIVE QA**
- 36.4 OPEN-DOMAIN QA**

CHAPTER 37: TEXT GENERATION

- 37.1 CHARACTER-LEVEL GENERATION**
- 37.2 LSTM-BASED GENERATION**
- 37.3 TRANSFORMER-BASED GENERATION**
- 37.4 CONTROLLED TEXT GENERATION**

CHAPTER 38: CONVERSATIONAL AI

- 38.1 CHATBOT FUNDAMENTALS**
- 38.2 INTENT + ENTITY SYSTEMS**
- 38.3 RETRIEVAL CHATBOTS**
- 38.4 GENERATIVE CONVERSATIONAL MODELS**

PART XI — SPEECH & MULTIMODAL NLP

CHAPTER 39: SPEECH AND AUDIO NLP

- 39.1 SPEECH RECOGNITION BASICS**
- 39.2 AUTOMATIC SPEECH RECOGNITION (ASR)**
- 39.3 TEXT-TO-SPEECH (TTS)**
- 39.4 SPEECH EMBEDDINGS**

CHAPTER 40: MULTIMODAL LANGUAGE SYSTEMS

- 40.1 VISION-LANGUAGE MODELS**
- 40.2 AUDIO-TEXT MODELS**
- 40.3 CROSS-MODAL RETRIEVAL**
- 40.4 MULTIMODAL FOUNDATION MODELS**

PART XII — PRODUCTION NLP SYSTEMS

CHAPTER 41: NLP MODEL DEPLOYMENT

- 41.1 FASTAPI DEPLOYMENT
- 41.2 STREAMLIT / GRADIO NLP APPS
- 41.3 DOCKERIZING NLP MODELS
- 41.4 BATCH VS REAL-TIME INFERENCE

CHAPTER 42: NLP OPTIMIZATION

- 42.1 QUANTIZATION
- 42.2 DISTILLATION
- 42.3 ONNX EXPORT
- 42.4 TENSORRT / RUNTIME OPTIMIZATION

CHAPTER 43: NLP MONITORING & MLOPS

- 43.1 DRIFT DETECTION FOR NLP
- 43.2 EMBEDDING DRIFT MONITORING
- 43.3 HALLUCINATION / OUTPUT MONITORING
- 43.4 RETRAINING PIPELINES

PART XIII — AZURE NLP & ENTERPRISE AI

CHAPTER 44: AZURE AI LANGUAGE SERVICES

- 44.1 AZURE AI LANGUAGE OVERVIEW
- 44.2 TEXT ANALYTICS
- 44.3 CUSTOM TEXT CLASSIFICATION
- 44.4 CONVERSATIONAL LANGUAGE UNDERSTANDING

CHAPTER 45: AZURE NLP DEPLOYMENT

- 45.1 AZURE ML FOR NLP MODELS
- 45.2 MANAGED ONLINE ENDPOINTS
- 45.3 BATCH ENDPOINTS
- 45.4 AZURE CONTAINER APPS

PART XIV — ENTERPRISE NLP PROJECTS

CHAPTER 46: BEGINNER PROJECTS

46.1 SPAM EMAIL CLASSIFIER

46.2 SENTIMENT ANALYZER

CHAPTER 47: INTERMEDIATE PROJECTS

47.1 NAMED ENTITY RECOGNITION SYSTEM

47.2 TEXT SUMMARIZATION PIPELINE

47.3 SEMANTIC SEARCH ENGINE

CHAPTER 48: ADVANCED PROJECTS

48.1 TRANSFORMER-BASED MACHINE TRANSLATION

48.2 QUESTION ANSWERING SYSTEM

48.3 ENTERPRISE CHATBOT PLATFORM

48.4 AZURE NLP PRODUCTION DEPLOYMENT

48.5 END-TO-END NLP MLOPS PIPELINE

APPENDICES

APPENDIX A — NLP INTERVIEW PREPARATION

APPENDIX B — HUGGING FACE DEEP DIVE

APPENDIX C — AZURE AI CERTIFICATION ROADMAP

APPENDIX D — RESEARCH PAPER READING FRAMEWORK
